

VISVESVARAYA TECHNOLOGICAL UNIVERSITY

"JnanaSangama", Belgaum -590014, Karnataka.



LAB

REPORT ON

MACHINE LEARNING

Submitted by

Imran Wadrli(1BM21CS077)

***in partial fulfillment for the award of the
degree of***

BACHELOR OF ENGINEERING

in

COMPUTER SCIENCE AND ENGINEERING



B. M. S. College of Engineering,

**Bull Temple Road, Bangalore
560019(March 2024 to June 2024)**



B. M. S. College of Engineering,

Bull Temple Road, Bangalore 560019

(Affiliated To Visveswaraya Technological University, Belgaum)

**Department of Computer Science and
Engineering**

CERTIFICATE

This is to certify that the Lab work entitled “**MACHINE LEARNING**” is carried out by **Imran Wadrali (1BM21CS077)** who is bonafide student of **B.M.S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visveswaraya Technological University, Belgaum during the year 2023-2024. The lab report has been approved as it satisfies the academic requirements in respect of **Machine Learning Lab - (22CS3PCMAL)** work prescribed for the said degree.

Dr. K. Panimozhi
Assistant Professor
BMSCE, Bengaluru

Dr. Jyothi S Nayak
Prof. & Head, Dept. of CSE
BMSCE, Bengaluru

Index		
Sl. No.	Experiment Title	Page No.
1	Write a python program to import and export data using Pandas library functions	4
2	Demonstrate various data pre-processing techniques for a given dataset	7
3	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	10
4	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	14
5	Build KNN Classification model for a given dataset.	17
6	Build Logistic Regression Model for a given dataset	19
7	Build Support vector machine model for a given dataset.	22
8	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	25
9	Implement Dimensionality reduction using Principle Component Analysis (PCA) method.	27
10	Build Artificial Neural Network model with back propagation on a given dataset	29
11	a) Implement Random forest ensemble method on a given dataset. b) Implement Boosting ensemble method on a given dataset.	32

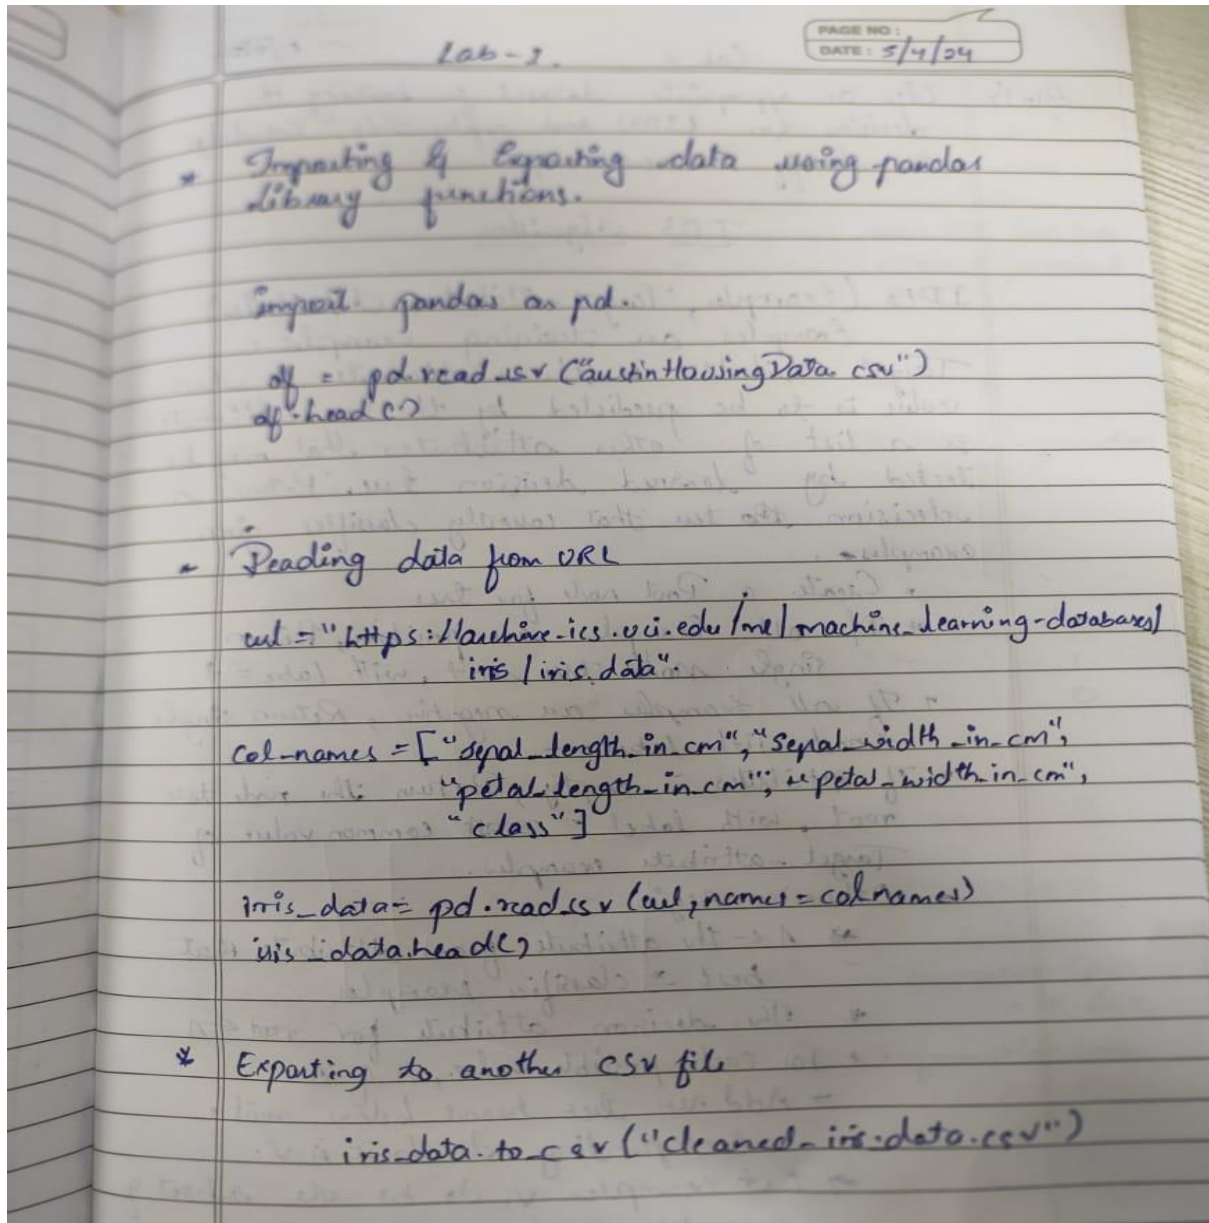
Course outcomes:

CO1	Apply machine learning techniques in computing systems
CO2	Evaluate the model using metrics
CO3	Design a model using machine learning to solve a problem
CO4	Conduct experiments to solve real-world problems using appropriate machine learning techniques

Lab 1

- 1) Write a python program to import and export data using Pandas library functions.

Algorithm (Observation book):



Code

```
import pandas as pd
df=pd.read_csv("/content/austinHousingData.csv")
df.head(5)
```

Output:

	zpid	city	streetAddress	zipcode	description	latitude	longitude	propertyTaxRate	garageSpaces	hasAssociation	...
0	111373431	pflugerville	14424 Lake Victor Dr	78660	14424 Lake Victor Dr, Pflugerville, TX 78660 I...	30.430632	-97.663078	1.98	2	True	...
1	120900430	pflugerville	1104 Strickling Dr	78660	Absolutely GORGEOUS 4 Bedroom home with 2 full...	30.432673	-97.661697	1.98	2	True	...
2	2084491383	pflugerville	1408 Fort Dessau Rd	78660	Under construction - estimated completion in A...	30.409748	-97.639771	1.98	0	True	...
3	120901374	pflugerville	1025 Strickling Dr	78660	Absolutely darling one story home in charming ...	30.432112	-97.661659	1.98	2	True	...
4	60134862	pflugerville	15005 Donna Jane Loop	78660	Brimming with appeal & warm livability! Sleek ...	30.437368	-97.656860	1.98	0	True	...

5 rows x 47 columns

- 2) Use an appropriate dataset for building the decision tree (ID3) and apply this knowledge to classify a new sample.

Algorithm (Observation book):

Lab 2

Q10) Use an appropriate dataset for building the decision tree (ID3) and apply this knowledge to classify a new sample.

ID3 Algorithm

ID3 (Examples, Target-attribute, Attributes)
Examples are training examples.
Target-attribute is the attribute whose value is to be predicted by the tree. Attributes is a list of other attributes that may be tested by learned decision tree. Returns a decision tree that correctly classifies given examples.

- + Create a Root node for tree
- * If all examples are positive, Return single node tree root, with label = +
- + If all examples are negative, Return single node with label = -
- * If Attributes is empty, Return the node tree root, with label = most common value of Target-attribute examples.
- + Otherwise Begin
 - * $A \leftarrow$ the attribute from Attributes that best classifies examples
 - * The decision attribute for root $\leftarrow A$
 - + For each possible value, v , of A
 - Add new tree branch below root, corresponding to best $A=v$.
 - Let examples v_i be the subset of examples with value v_i for A

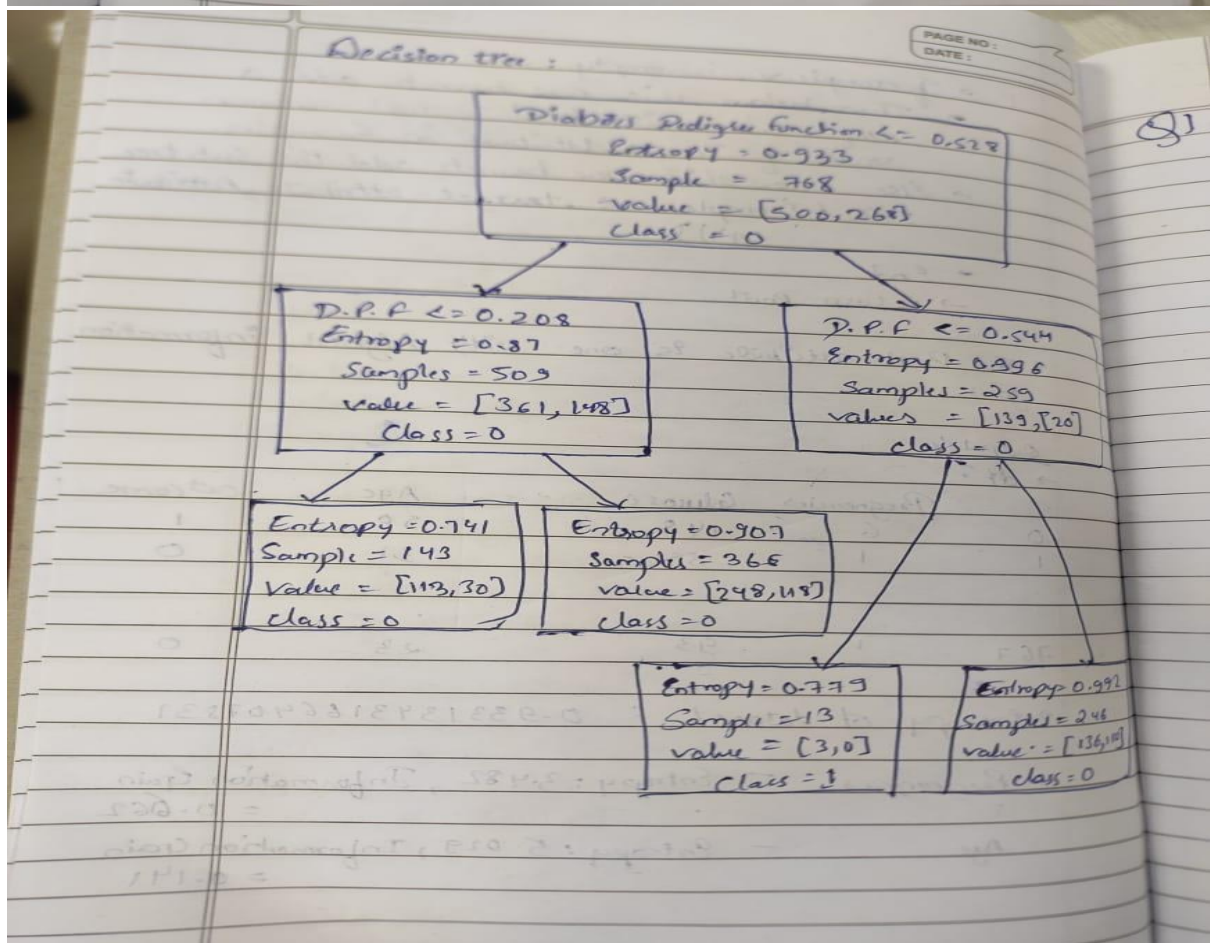
= If example v_i is empty
 = Then below this new branch add a leaf node with label = most common value of Target attribute in Examples.
 → Else below this new branch add the subtree
 103. (Examples v_i : target-attribute, Attribute (A4))
 = End
 → Return Root

* The Best attribute is one with highest information gain.

Output:

	Pregnancies	Glucose	Age	Outcome
0	6	148	50	1
1	1	85	31	0
767	1	93	23	0

→ Entropy of dataset : 0.9331343166407831
 → Pregnancies : Entropy : 3.482, Information Gain = 0.662
 Age : Entropy : 5.029, Information Gain = 0.141



Code:

```
import pandas as pd
from sklearn.tree import DecisionTreeClassifier, plot_tree
import matplotlib.pyplot as plt
import math

df = pd.read_csv('/content/diabetes.csv')

def calculate_entropy(data, target_column):
    total_rows = len(data)
    target_values = data[target_column].unique()

    entropy = 0
    for value in target_values:
        # Calculate the proportion of instances with the current value
        value_count = len(data[data[target_column] == value])
        proportion = value_count / total_rows
        entropy -= proportion * math.log2(proportion)

    return entropy

entropy_outcome = calculate_entropy(df, 'Outcome')
print(f"Entropy of the dataset: {entropy_outcome}")

def calculate_entropy(data, target_column): # for each categorical variable
    total_rows = len(data)
    target_values = data[target_column].unique()

    entropy = 0
    for value in target_values:
        # Calculate the proportion of instances with the current value
        value_count = len(data[data[target_column] == value])
        proportion = value_count / total_rows
        entropy -= proportion * math.log2(proportion) if proportion != 0 else 0

    return entropy

def calculate_information_gain(data, feature, target_column):

    # Calculate weighted average entropy for the feature
    unique_values = data[feature].unique()
    weighted_entropy = 0

    for value in unique_values:
        subset = data[data[feature] == value]
        proportion = len(subset) / len(data)
        weighted_entropy += proportion * calculate_entropy(subset, target_column)
```



```

# Calculate information gain
information_gain = entropy_outcome - weighted_entropy

return information_gain

for column in df.columns[:-1]:
    entropy = calculate_entropy(df, column)
    information_gain = calculate_information_gain(df, column, 'Outcome')
    print(f"{column} - Entropy: {entropy:.3f}, Information Gain: {information_gain:.3f}")

# Feature selection for the first step in making decision tree
selected_feature = 'DiabetesPedigreeFunction'

# Create a decision tree
clf = DecisionTreeClassifier(criterion='entropy', max_depth=1)
X = df[[selected_feature]]
y = df['Outcome']
clf.fit(X, y)

plt.figure(figsize=(8, 6))
plot_tree(clf, feature_names=[selected_feature], class_names=['0', '1'], filled=True,
rounded=True)
plt.show()

def id3(data, target_column, features):
    if len(data[target_column].unique()) == 1:
        return data[target_column].iloc[0]

    if len(features) == 0:
        return data[target_column].mode().iloc[0]

    best_feature = max(features, key=lambda x: calculate_information_gain(data, x,
target_column))

    tree = {best_feature: {}}

    features = [f for f in features if f != best_feature]

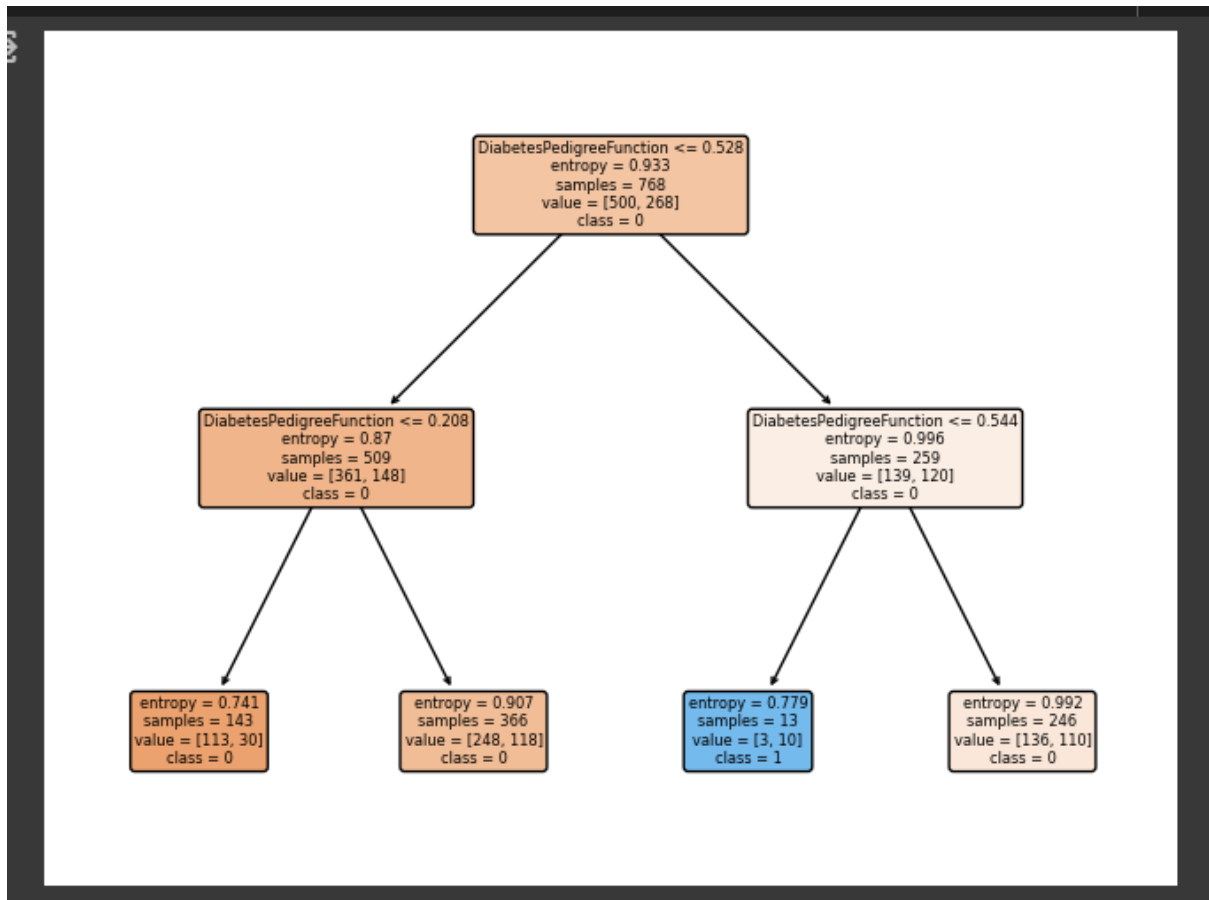
    for value in data[best_feature].unique():
        subset = data[data[best_feature] == value]
        tree[best_feature][value] = id3(subset, target_column, features)

    return tree

id3(df, 'Outcome', ['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness', 'Insulin',
'BMI', 'DiabetesPedigreeFunction', 'Age'] )

```

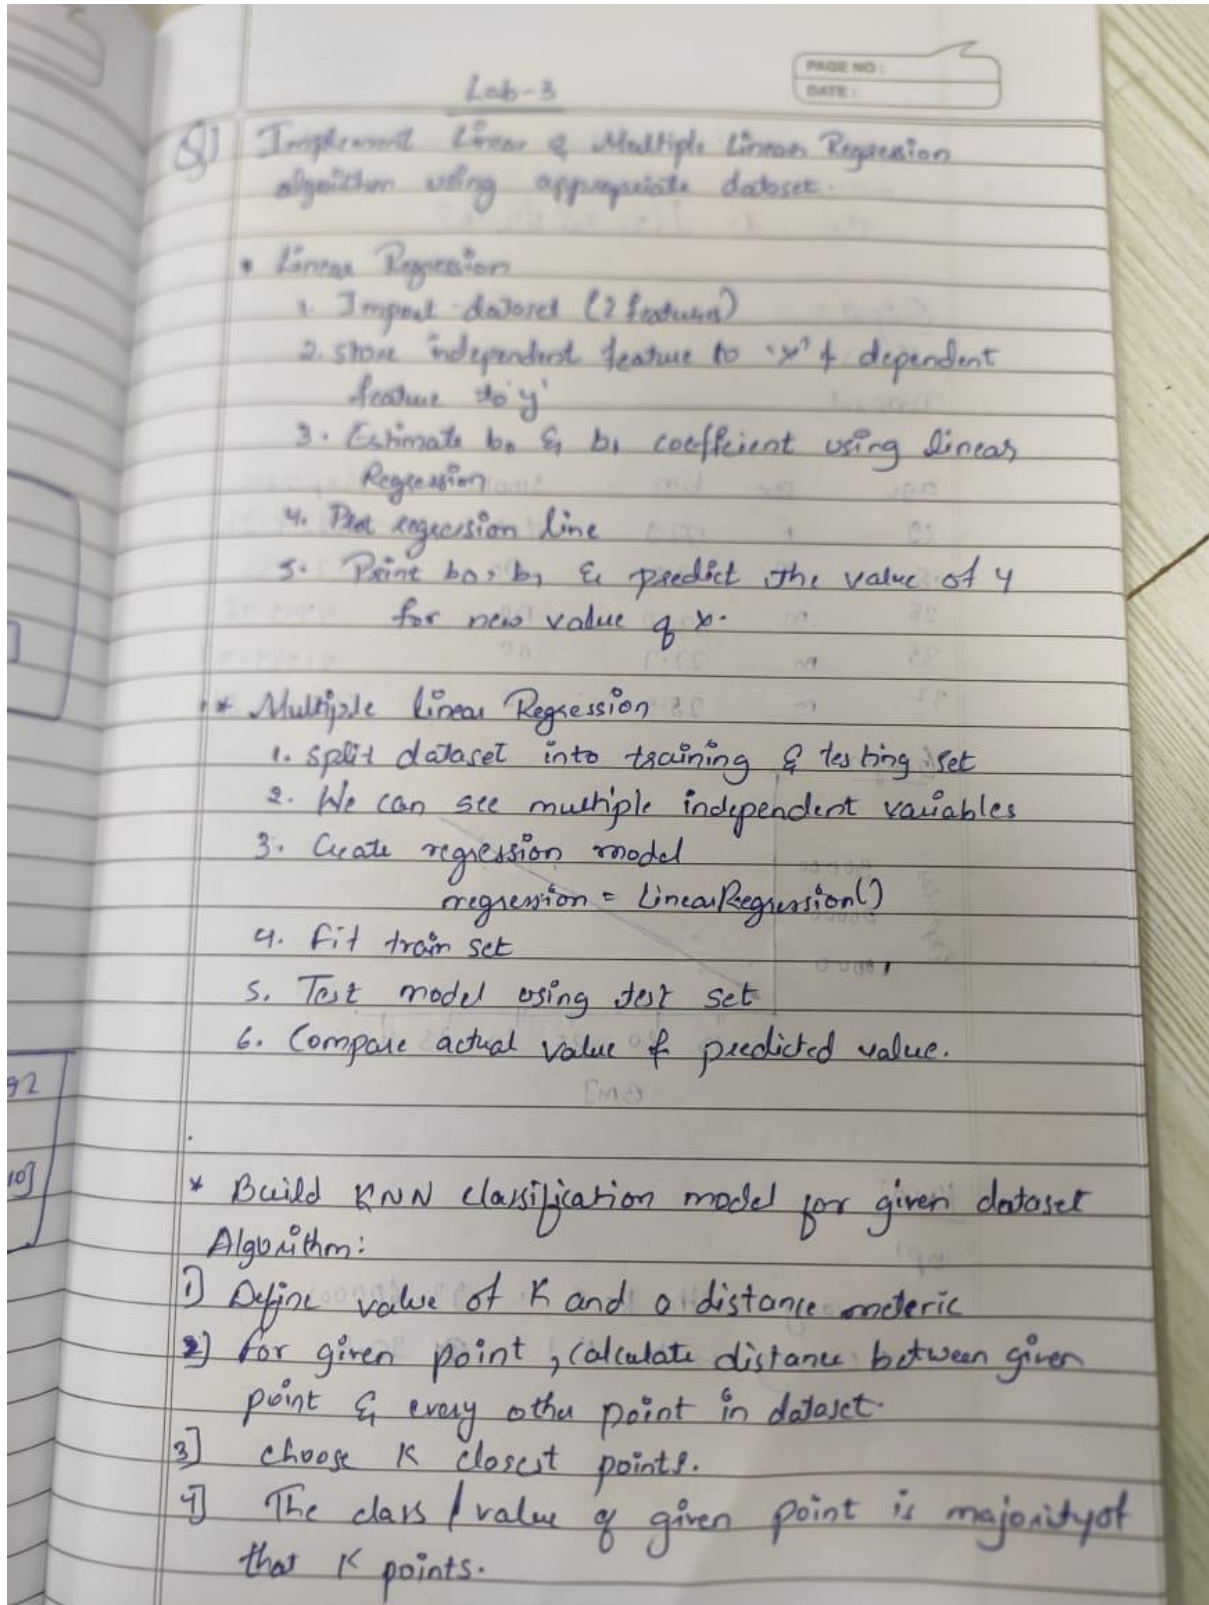
	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1



Pregnancies - Entropy: 3.482, Information Gain: 0.062
 Glucose - Entropy: 6.751, Information Gain: 0.304
 BloodPressure - Entropy: 4.792, Information Gain: 0.059
 SkinThickness - Entropy: 4.586, Information Gain: 0.082
 Insulin - Entropy: 4.682, Information Gain: 0.277
 BMI - Entropy: 7.594, Information Gain: 0.344
 DiabetesPedigreeFunction - Entropy: 8.829, Information Gain: 0.651
 Age - Entropy: 5.029, Information Gain: 0.141

3) Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

Algorithm



Code:

```
import numpy as np
import matplotlib.pyplot as plt

[ ] def estimate_coef(x, y):
    # number of observations/points
    n = np.size(x)

    # mean of x and y vector
    m_x = np.mean(x)
    m_y = np.mean(y)

    # calculating cross-deviation and deviation about x
    SS_xy = np.sum(y*x) - n*m_y*m_x
    SS_xx = np.sum(x*x) - n*m_x*m_x

    # calculating regression coefficients
    b_1 = SS_xy / SS_xx
    b_0 = m_y - b_1*m_x

    return (b_0, b_1)
```

```
def plot_regression_line(x, y, b):
    # plotting the actual points as scatter plot
    plt.scatter(x, y, color = "m",
               marker = "o", s = 30)

    # predicted response vector
    y_pred = b[0] + b[1]*x

    # plotting the regression line
    plt.plot(x, y_pred, color = "g")

    # putting labels
    plt.xlabel('x')
    plt.ylabel('y')
```


```
[ ] import pandas as pd
def main():
    # observations / data
    df=pd.read_csv("Salary_dataset.csv")
    x = np.array(df['YearsExperience'])
    y = np.array(df['Salary'])

    # estimating coefficients
    b = estimate_coef(x, y)
    plot_regression_line(x,y,b)
    print("Estimated coefficients:\nb_0 = {}\nb_1 = {}".format(b[0],b[1]))
    x=int(input("Enter X value:"))
    y=b[0]+b[1]*x
    print(f"The predicted salary is:{y}")
```

```
[ ] main()
```

OUTPUT:

```
Estimated coefficients:
b_0 = 24848.203966523113
b_1 = 9449.962321455092
Enter X value45
The predicted salary is:450096.5084320023
```



b) Implement Multi-Linear Regression algorithm using appropriate dataset

```
[6] """
Standardizes the input data using mean and standard deviation.

Parameters:
    X_train (numpy.ndarray): Training data.
    X_test (numpy.ndarray): Testing data.

Returns:
    Tuple of standardized training and testing data.
"""
# Calculate the mean and standard deviation using the training data
mean = np.mean(X_train, axis=0)
std = np.std(X_train, axis=0)

# Standardize the data
X_train = (X_train - mean) / std
X_test = (X_test - mean) / std

return X_train, X_test

X_train, X_test = standardize_data(X_train, X_test)

[7] X_train = np.expand_dims(X_train, axis=-1)
X_test = np.expand_dims(X_test, axis=-1)

[8] class LinearRegression:
    """
    Linear Regression Model with Gradient Descent

    Linear regression is a supervised machine learning algorithm used for modeling the relationship
    between a dependent variable (target) and one or more independent variables (features) by fitting
    a linear equation to the observed data.

    This class implements a linear regression model using gradient descent optimization for training.
    It provides methods for model initialization, training, prediction, and model persistence.

    Parameters:
        learning_rate (float): The learning rate used in gradient descent.
        convergence_tol (float, optional): The tolerance for convergence (stopping criterion). Defaults to 1e-6.

    Attributes:
        W (numpy.ndarray): Coefficients (weights) for the linear regression model.
        b (float): Intercept (bias) for the linear regression model.

    Methods:
        initialize_parameters(n_features): Initialize model parameters.
        forward(X): Compute the forward pass of the linear regression model.
        compute_cost(predictions): Compute the mean squared error cost.
        backward(predictions): Compute gradients for model parameters.
        fit(X, y, iterations, plot_cost=True): Fit the linear regression model to training data.
        predict(X): Predict target values for new input data.
        save_model(filename=None): Save the trained model to a file using pickle.
        load_model(filename): Load a trained model from a file using pickle.

    Examples:
        >>> from linear_regression import LinearRegression
        >>> model = LinearRegression(learning_rate=0.01)
        >>> model.fit(X_train, y_train, iterations=1000)
        >>> predictions = model.predict(X_test)
        """

    def __init__(self, learning_rate, convergence_tol=1e-6):
        self.learning_rate = learning_rate
        self.convergence_tol = convergence_tol
        self.W = None
        self.b = None

    def initialize_parameters(self, n_features):
        """
        Initialize model parameters.

        Parameters:
            n_features (int): The number of features in the input data.
        """
        self.W = np.random.randn(n_features) * 0.01
        self.b = 0

    def forward(self, X):
        """
        Compute the forward pass of the linear regression model.

        Parameters:
            X (numpy.ndarray): Input data of shape (m, n_features).

```



```
[8] Returns:
      numpy.ndarray: Predictions of shape (m,).
      """
      return np.dot(X, self.W) + self.b

def compute_cost(self, predictions):
    """
    Compute the mean squared error cost.

    Parameters:
        predictions (numpy.ndarray): Predictions of shape (m,).

    Returns:
        float: Mean squared error cost.
    """
    m = len(predictions)
    cost = np.sum(np.square(predictions - self.y)) / (2 * m)
    return cost

def backward(self, predictions):
    """
    Compute gradients for model parameters.

    Parameters:
        predictions (numpy.ndarray): Predictions of shape (m,).

    Updates:
        numpy.ndarray: Gradient of W.
        float: Gradient of b.
    """
    m = len(predictions)
    self.dW = np.dot(predictions - self.y, self.X) / m
    self.db = np.sum(predictions - self.y) / m
def fit(self, X, y, iterations, plot_cost=True):
    """
    Fit the linear regression model to the training data.

    Parameters:
        X (numpy.ndarray): Training input data of shape (m, n_features).
        y (numpy.ndarray): Training labels of shape (m,).
        iterations (int): The number of iterations for gradient descent.
        plot_cost (bool, optional): Whether to plot the cost during training. Defaults to True.
```

```
[8] Raises:
      AssertionError: If input data and labels are not NumPy arrays or have mismatched shapes.

Plots:
    Plotly line chart showing cost vs. iteration (if plot_cost is True).
    """
    assert isinstance(X, np.ndarray), "X must be a NumPy array"
    assert isinstance(y, np.ndarray), "y must be a NumPy array"
    assert X.shape[0] == y.shape[0], "X and y must have the same number of samples"
    assert iterations > 0, "Iterations must be greater than 0"

    self.X = X
    self.y = y
    self.initialize_parameters(X.shape[1])
    costs = []

    for i in range(iterations):
        predictions = self.forward(X)
        cost = self.compute_cost(predictions)
        self.backward(predictions)
        self.W -= self.learning_rate * self.dW
        self.b -= self.learning_rate * self.db
        costs.append(cost)

        if i % 100 == 0:
            print(f'Iteration: {i}, Cost: {cost}')

        if i > 0 and abs(costs[-1] - costs[-2]) < self.convergence_tol:
            print(f'Converged after {i} iterations.')
            break

    if plot_cost:
        fig = px.line(y=costs, title="Cost vs Iteration", template="plotly_dark")
        fig.update_layout(
            title_font_color="#41BEE9",
            xaxis=dict(color="#41BEE9", title="Iterations"),
            yaxis=dict(color="#41BEE9", title="Cost")
        )

    fig.show()
```

```

fig.show()
def predict(self, X):
    """
    Predict target values for new input data.

    Parameters:
        X (numpy.ndarray): Input data of shape (m, n_features).

    Returns:
        numpy.ndarray: Predicted target values of shape (m,).
    """
    return self.forward(X)

def save_model(self, filename=None):
    """
    Save the trained model to a file using pickle.

    Parameters:
        filename (str): The name of the file to save the model to.
    """
    model_data = {
        'learning_rate': self.learning_rate,
        'convergence_tol': self.convergence_tol,
        'W': self.W,
        'b': self.b
    }

    with open(filename, 'wb') as file:
        pickle.dump(model_data, file)

@classmethod
def load_model(cls, filename):
    """
    Load a trained model from a file using pickle.

    Parameters:
        filename (str): The name of the file to load the model from.

```

```

8] with open(filename, 'wb') as file:
    pickle.dump(model_data, file)

@classmethod
def load_model(cls, filename):
    """
    Load a trained model from a file using pickle.

    Parameters:
        filename (str): The name of the file to load the model from.

    Returns:
        LinearRegression: An instance of the LinearRegression class with loaded parameters.
    """
    with open(filename, 'rb') as file:
        model_data = pickle.load(file)

    # Create a new instance of the class and initialize it with the loaded parameters
    loaded_model = cls(model_data['learning_rate'], model_data['convergence_tol'])
    loaded_model.W = model_data['W']
    loaded_model.b = model_data['b']

    return loaded_model

```

```

9] lr = LinearRegression(0.01)
    lr.fit(X_train, y_train, 10000)

```

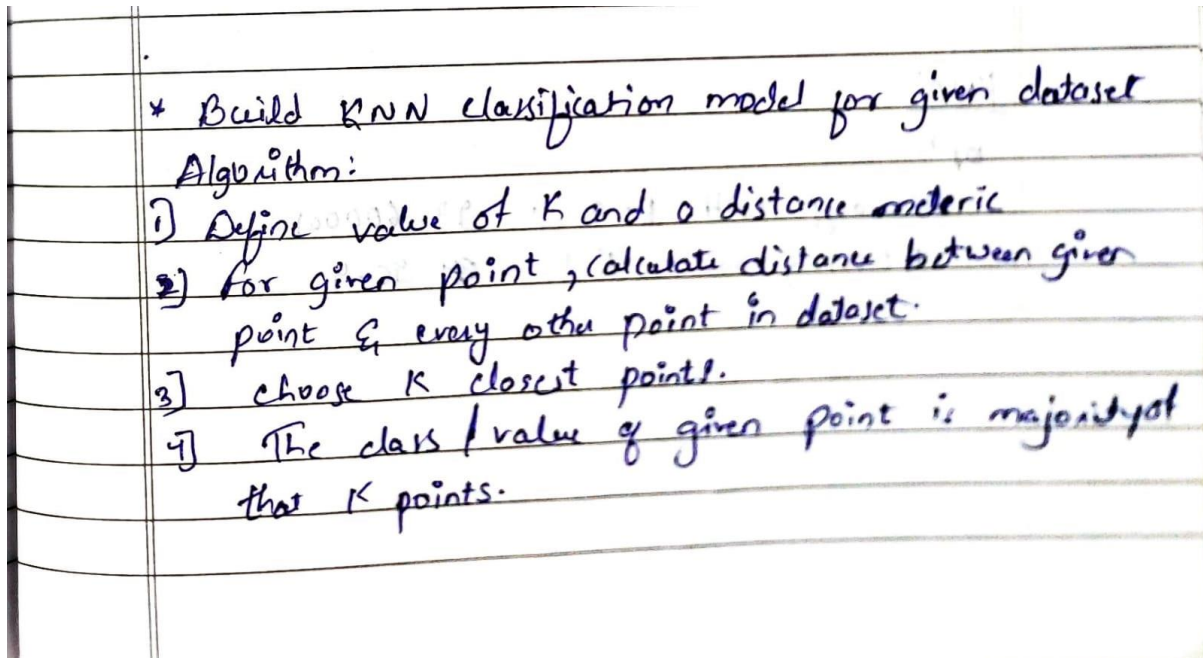
```

Iteration: 0, Cost: 1670.0184887161677
Iteration: 100, Cost: 227.15535101517312
Iteration: 200, Cost: 33.84101696145528
Iteration: 300, Cost: 7.9408253395546575
Iteration: 400, Cost: 4.4707260872934835
Iteration: 500, Cost: 4.005803317750673
Iteration: 600, Cost: 3.943513116253261
Iteration: 700, Cost: 3.9351674953098015
Iteration: 800, Cost: 3.9340493517293096
Converged after 863 iterations.

```

4) Build KNN Classification model for a given dataset.

Algorithm



Code:

```
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix
from sklearn import datasets
iris = datasets.load_iris()

x = iris.data
y = iris.target

print('sepal-length', 'sepal-width', 'petal-length', 'petal-width')
print(x)
print('class: 0 - Iris-Setosa, 1 - Iris-Versicolour, 2 - Iris-Virginica')
print(y)
```

```

In [5]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3)

#To Training the model and Nearest nighbors K=5
classifier = KNeighborsClassifier(n_neighbors=5)
classifier.fit(x_train, y_train)

#to make predictions on our test data
y_pred=classifier.predict(x_test)

print('Prediction -')

for i,test in enumerate(x_test) :
    print(f'{test} - {y_pred[i]}')

# print('Confusion Matrix')
# print(confusion_matrix(y_test,y_pred))
# print('Accuracy Metrics')
# print(classification_report(y_test,y_pred))

```

Results:

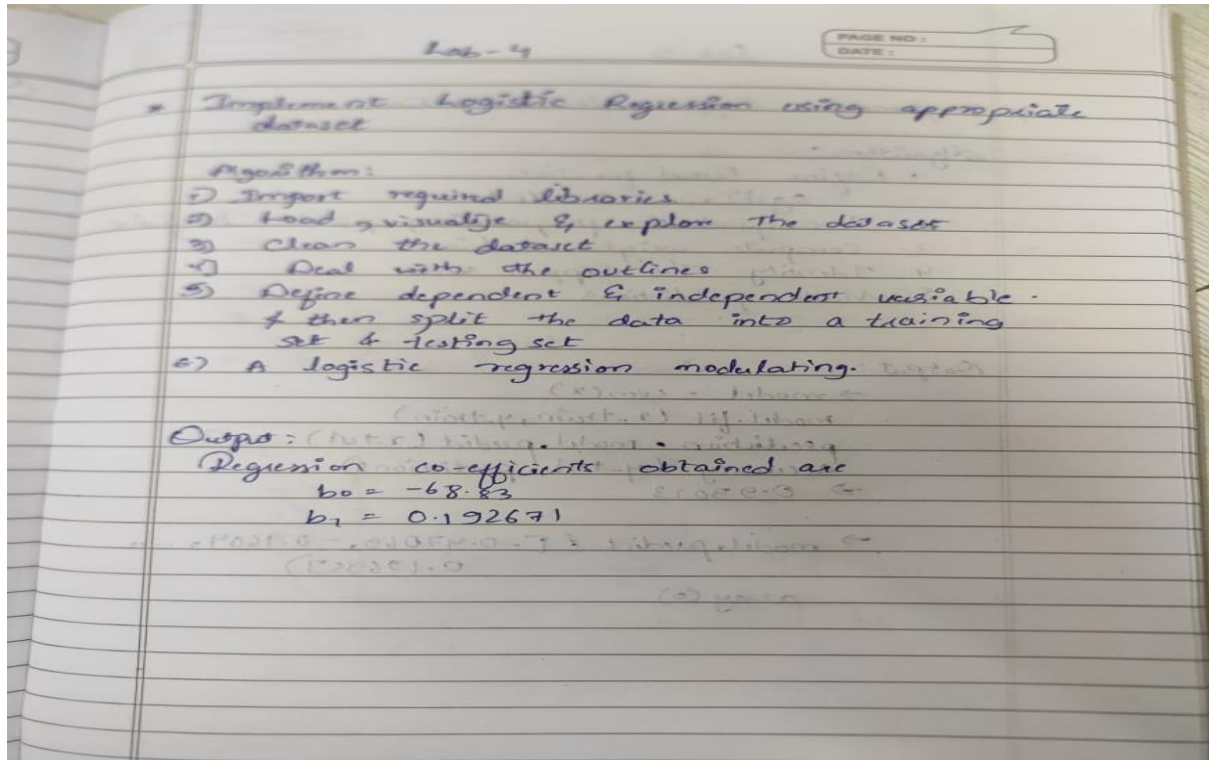
```

Prediction -
[5.2 4.1 1.5 0.1] - 0
[5.5 2.3 4.  1.3] - 1
[6.7 3.1 4.7 1.5] - 1
[7.  3.2 4.7 1.4] - 1
[6.2 2.8 4.8 1.8] - 2
[5.7 2.8 4.5 1.3] - 1
[6.  3.4 4.5 1.6] - 1
[5.1 3.8 1.6 0.2] - 0
[5.5 2.5 4.  1.3] - 1
[4.8 3.1 1.6 0.2] - 0
[6.1 3.  4.9 1.8] - 2
[4.7 3.2 1.6 0.2] - 0
[5.6 2.9 3.6 1.3] - 1
[5.4 3.9 1.3 0.4] - 0
[5.  3.2 1.2 0.2] - 0
[6.1 2.9 4.7 1.4] - 1
[5.  3.4 1.5 0.2] - 0
[7.7 2.8 6.7 2.  ] - 2
[4.6 3.2 1.4 0.2] - 0
[5.7 2.9 4.2 1.3] - 1
[4.6 3.6 1.  0.2] - 0
[6.8 2.8 4.8 1.4] - 1
[6.8 3.2 5.9 2.3] - 2

```

5) Build Logistic Regression Model for a given dataset

Algorithm:



Code:



```

age
12  27
13  29
4   46
23  45
3   52
8   62
1   25
11  28
25  54

```

```

In [32]: from sklearn.linear_model import LogisticRegression
model = LogisticRegression()
model.fit(X_train, y_train)
print(X_test)
y_predicted = model.predict(X_test)
probab = model.predict_proba(X_test)
score = model.score(X_test,y_test)
print(y_predicted)
print("\nprobability: ")
print(probab)
print("\naccuracy: ")
print(score)

```

```

age
12  27
13  29
4   46
23  45
3   52
8   62
1   25
11  28
25  54
[0 0 1 1 1 1 0 0 1]

```

```

probability:
[[0.75147045 0.24852955]
 [0.69990447 0.30009553]
 [0.20424226 0.79575774]
 [0.2261509  0.7738491 ]
 [0.10537592 0.89462408]
 [0.03115876 0.96884124]
 [0.79674889 0.20325111]
 [0.7264443  0.2735557 ]
 [0.08328741 0.91671259]]

```

```

accuracy:
0.8888888888888888

```

```

In [33]: print(X_test)

```

```

age
12  27
13  29
4   46
23  45
3   52
8   62
1   25
11  28
25  54

```

```

In [38]: import math
def sigmoid(x):
    return 1 / (1 + math.exp(-x))

```

```

In [39]: def prediction_function(age):
z = 0.042 * age - 1.53 # 0.04150133 ~ 0.042 and -1.52726963 ~ -1.53
y = sigmoid(z)
return y

```



```
In [40]: def insure(probability):  
        if (probability > 0.5):  
            print("The customer will get insurance.")  
        else:  
            print("The customer will not get insurance.")  
  
        age = 35  
        probability = prediction_function(age)  
        print("Probability of buying insurance for age 35:", probability)  
        insure(probability)
```

Probability of buying insurance for age 35: 0.4850044983805899
The customer will not get insurance.

```
In [41]: age = 43  
        probability = prediction_function(age)  
        print("Probability of buying insurance for age 43:", probability)  
        insure(probability)
```

Probability of buying insurance for age 43: 0.568565299077705
The customer will get insurance.

Results:

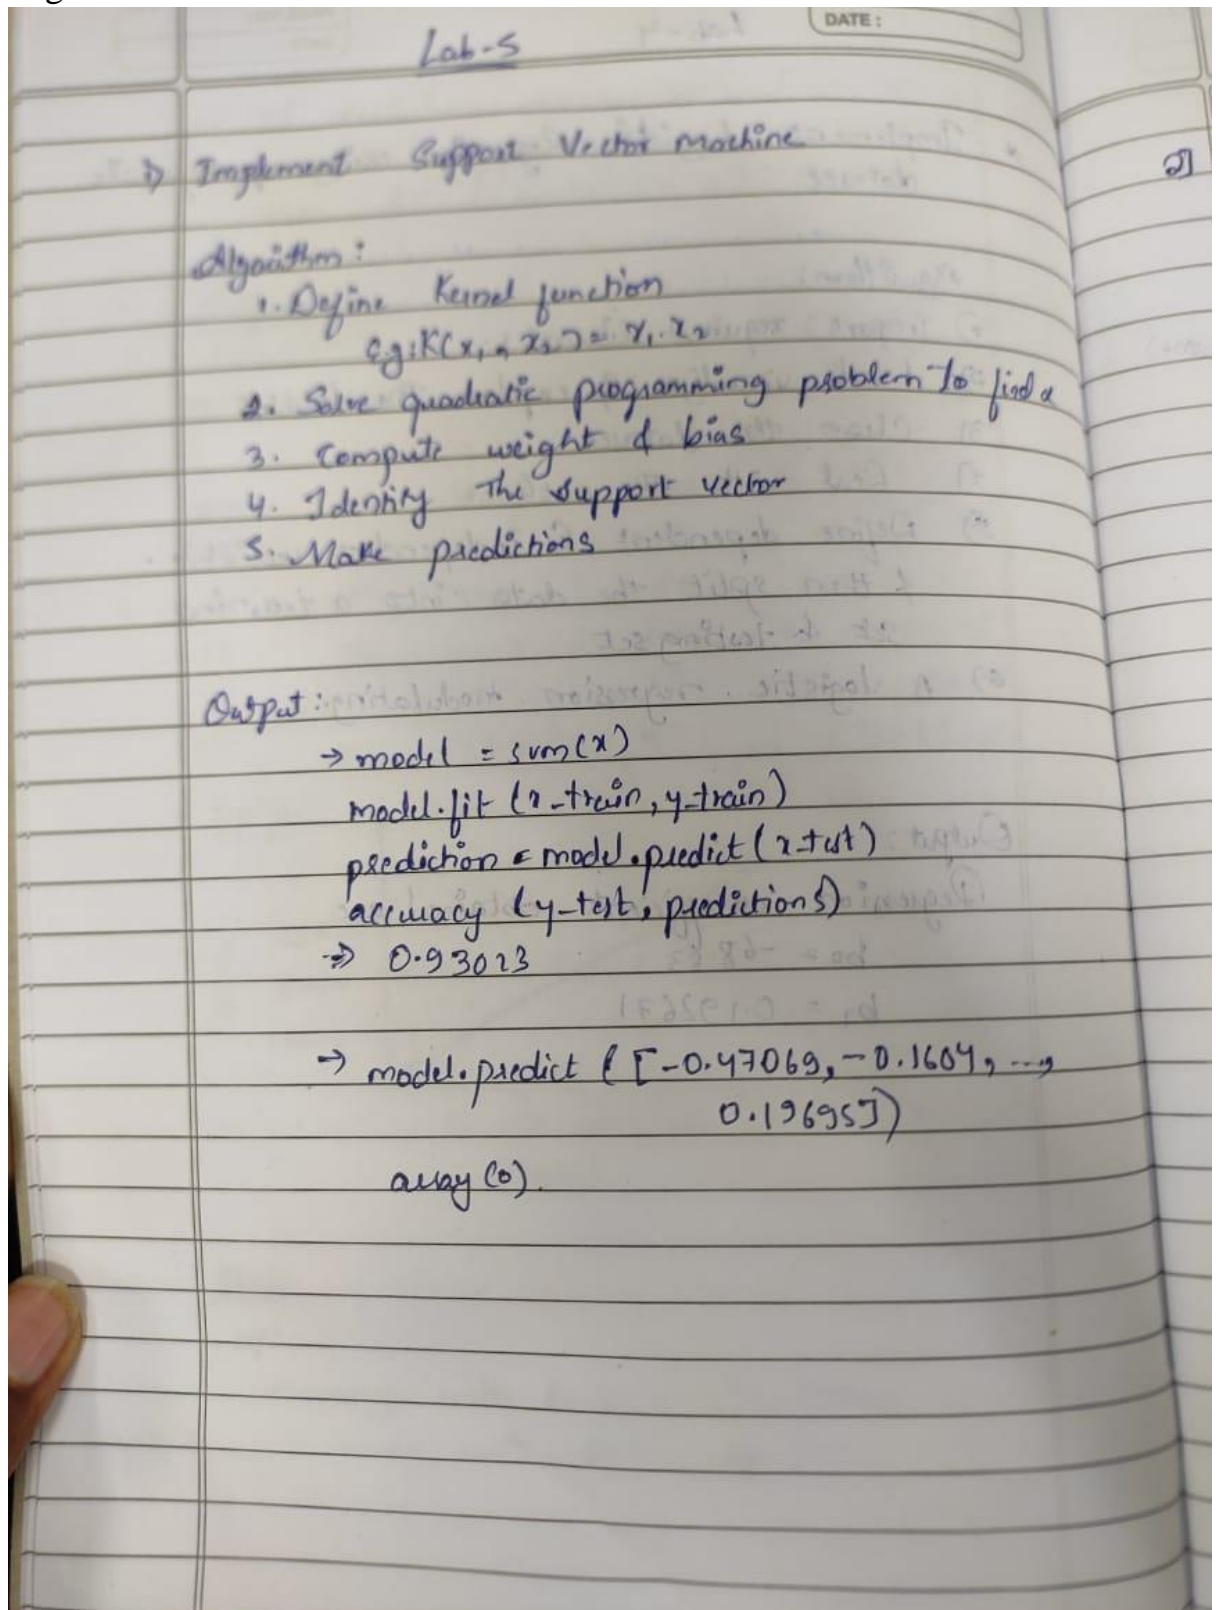
Probability of buying insurance for age 35: 0.4850044983805899
The customer will not get insurance.

```
In [41]: age = 43  
        probability = prediction_function(age)  
        print("Probability of buying insurance for age 43:", probability)  
        insure(probability)
```

Probability of buying insurance for age 43: 0.568565299077705
The customer will get insurance.

6) Build Support vector machine model for a given dataset

Algorithm:



Code:

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score
```

```
In [2]: # Load the Iris dataset
iris = load_iris()
```

```
In [3]: # Convert the dataset into a pandas DataFrame
iris_df = pd.DataFrame(data=iris.data, columns=iris.feature_names)
iris_df['target'] = iris.target
```


```
In [4]: # Display the first few rows of the DataFrame
print(iris_df.head())
```

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

	target
0	0
1	0
2	0
3	0
4	0

In [5]:

```
# Plotting to visualize the data
plt.figure(figsize=(10, 6))
plt.scatter(iris_df['sepal length (cm)'], iris_df['sepal width (cm)'],
            c=iris_df['target'], cmap='viridis')
plt.xlabel('Sepal Length (cm)')
plt.ylabel('Sepal Width (cm)')
plt.title('Iris Dataset - Sepal Length vs Sepal Width')
plt.colorbar(label='Species')
plt.show()
```



In [6]:

```
# Splitting the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(iris.data, iris.target, test_size=0.3, random_state=42)
```

In [7]:

```
# Creating and training the SVM classifier
svm_classifier = SVC(kernel='linear')
svm_classifier.fit(X_train, y_train)

# Predicting the Labels for the test set
y_pred = svm_classifier.predict(X_test)

# Calculating the accuracy of the model
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy of SVM Classifier:", accuracy)
```

Results:

Accuracy of SVM Classifier: 1.0

In [8]: y_pred

Out[8]: array([1, 0, 2, 1, 1, 0, 1, 2, 1, 1, 2, 0, 0, 0, 1, 2, 1, 1, 2, 0, 2,
0, 2, 2, 2, 2, 0, 0, 0, 0, 1, 0, 0, 2, 1, 0, 0, 0, 2, 1, 1, 0,
0])

7) Build k-Means algorithm to cluster a set of data stored in a .CSV file.

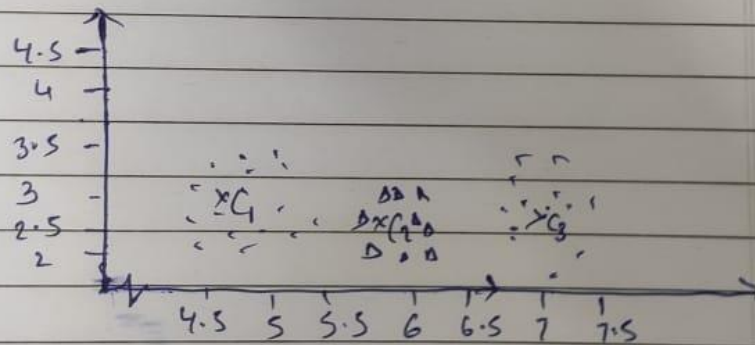
Algorithm:

Q) K-means clustering

Algorithm:

1. Select number of K to decide no. of clusters.
2. Select random ' k ' points [Centroids]
3. Assign each point to closest centroid, which will form the predefined ' k ' clusters
4. Calculate variance and place a new centroid of each cluster
5. Repeat step 3, reassign the centroid.
6. If any reassignment go to step (4) else go to finish
7. The model is trained

Output:



Code:

```
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.cluster import KMeans
import pandas as pd
import numpy as np

# Import some data to play with
iris = datasets.load_iris()
X = pd.DataFrame(iris.data)
X.columns = ['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width']
y = pd.DataFrame(iris.target)
y.columns = ['Targets']

# Build the K Means Model
model = KMeans(n_clusters=3)
model.fit(X) # model.labels_ gives cluster no for which samples belongs

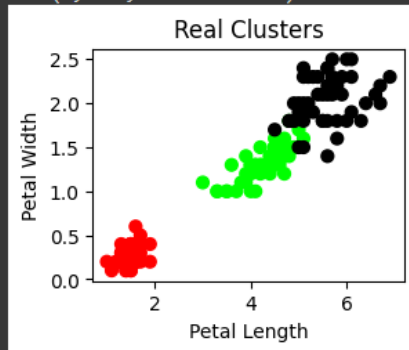
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of "n_init" will change from 10 to 'auto' in 1.4. Set the value of "n_init" explicitly to suppress the warning
  warnings.warn(
  KMeans
  KMeans(n_clusters=3)

# Visualise the clustering results
plt.figure(figsize=(14,14))
colormap = np.array(['red', 'lime', 'black'])

<Figure size 1408x1408 with 0 Axes>


# Plot the Original Classifications using Petal features
plt.subplot(2, 2, 1)
plt.scatter(X.Petal_Length, X.Petal_Width, c=colormap[y.Targets], s=40)
plt.title('Real Clusters')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')
```

Text(0, 0.5, 'Petal Width')



```
# Plot the Models Classifications
plt.subplot(2, 2, 2)
plt.scatter(X.Petal_Length, X.Petal_Width, c=colormap[model.labels_], s=40)
plt.title('K-Means Clustering')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')
```

Text(0, 0.5, 'Petal Width')



8) Implement Dimensionality reduction using Principle Component Analysis (PCA) method.

Algorithm:

3) Principal component Analysis:

Algorithm:

- 1) Calculate mean
- 2) Calculation of Covariance matrix
- 3) Eigen values of covariance matrix
- 4) Computation of eigen-vector (Unit-eigen vector)
- 5) Computation of first principal components
- 6) Geometric meaning of first principal component

Output:

pca.explained_variance_ratio
array ([0.9839448, 0.01620498])

Code:

```
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
import pandas as pd
import numpy as np

# Load the iris dataset
iris = datasets.load_iris()
X = pd.DataFrame(iris.data, columns=['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width'])
y = pd.DataFrame(iris.target, columns=['Targets'])

# Standardize the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Apply PCA
pca = PCA(n_components=2)
X_pca = pca.fit_transform(X_scaled)

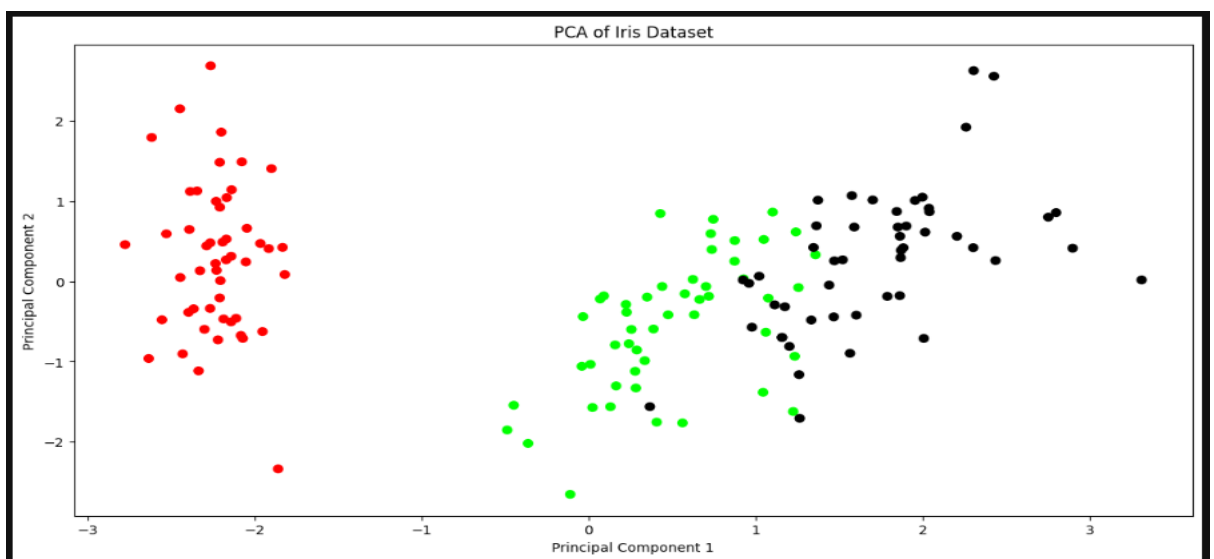
# Convert PCA result to a DataFrame
X_pca_df = pd.DataFrame(X_pca, columns=['PCA1', 'PCA2'])

# Add the target column for visualization
X_pca_df['Targets'] = y.Targets

# Visualize the PCA result
plt.figure(figsize=(14, 7))
colormap = np.array(['red', 'lime', 'black'])

# Plot the PCA transformed data
plt.scatter(X_pca_df.PCA1, X_pca_df.PCA2, c=colormap[X_pca_df.Targets], s=40)
plt.title('PCA of Iris Dataset')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.show()
```

Results:



9) Build Artificial Neural Network model with back propagation on a given dataset

Algorithm:

PAGE NO:
 DATE:

Lab-6

* Artificial Neural network with back propagation

Algorithm

- 1) * Initialize Parameters
 - + Normalise i/p features matrix 'n'
 - normalise o/p y
 - Set hyper parameters: 'no-of-epochs', 'no of neurons'
- 2) Define activation func
 - Sigmoid function adjustments
- 3) Training network
 - Forward propagation
 - + compute i/p to hidden layer
 - + Add bias
 - + Apply Activation func
- 4) Backward propagation
 - + compute error
 - + compute gradient
 - + compute delta
- 5) Update weights & biases.

Old wts:
$$\begin{bmatrix} 0.6667 & 2 \\ 0.333 & 0.556 \end{bmatrix}$$

New wts:
$$\begin{bmatrix} 0.97 & 0.1 & 0.667 \end{bmatrix}$$

predicted o/p:
$$\begin{bmatrix} 0.80056911 & 0.79392911 \\ 0.85112342 \end{bmatrix}$$

Code:

```

In [10]: import numpy as np
x = np.array([[2,9],[1,5],[3,6]],dtype = float)
y = np.array([[92],[86],[89]],dtype = float)

x = x/np.amax(x,axis=0)
y = y/100

In [11]: #Variable Initialization
epoch = 5000
lr = 0.1
inputlayer_neurons = 2
hiddenlayer_neurons = 3
output_neurons = 1

In [12]: # weight and bias Initialization

wh = np.random.uniform(size=(inputlayer_neurons,hiddenlayer_neurons))
bh = np.random.uniform(size=(1,hiddenlayer_neurons))
wout = np.random.uniform(size=(hiddenlayer_neurons,output_neurons))
bout = np.random.uniform(size=(1,output_neurons))

In [13]: #sigmoid function

def sigmoid(x):
    return 1/(1+np.exp(-x))

# Derivative of Sigmoid

def der_sigmoid(x):
    return x*(1-x)

```

```

In [14]: # Draws a random range of numbers uniformly of dim x*y

for i in range(epoch):

    # forward propagation
    hinp1 = np.dot(x,wh)
    hinp = hinp1 + bh
    hlayer_act = sigmoid(hinp)
    outinp1 = np.dot(hlayer_act,wout)
    outinp = outinp1 + bout
    output = sigmoid(outinp)

    # Backpropagation
    EO = y - output
    outgrad = der_sigmoid(output)
    d_output = EO*outgrad
    EH = d_output.dot(wout.T)

In [15]: # how much hidden Layer weights contributed to error

    hiddengrad = der_sigmoid(hlayer_act)
    d_hiddenlayer = EH*hiddengrad

    #dotproduct of nextlayererror and current layer op

    wout += hlayer_act.T.dot(d_output)*lr
    wh += x.T.dot(d_hiddenlayer)*lr

    print("Input: \n" + str(x))
    print("Actual output: \n" + str(y))
    print("Predicted Output: \n",output)

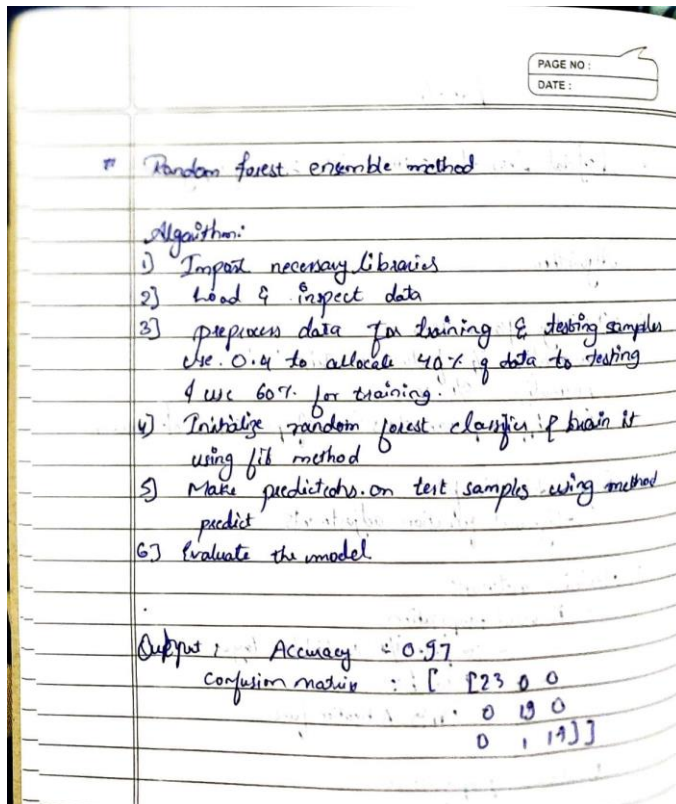
```

Results

```
Input:
[[0.66666667 1.      ]
 [0.33333333 0.55555556]
 [1.         0.66666667]]
Actual output:
[[0.92]
 [0.86]
 [0.89]]
Predicted Output:
[[0.80056875]
 [0.79393831]
 [0.80112347]]
```

10) Implement Random forest ensemble method on a given dataset.

Algorithm:



Code:

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn import datasets

# Load the data
iris_data = datasets.load_iris()

X = pd.DataFrame(iris_data.data, columns=['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width'])
y = pd.DataFrame(iris_data.target, columns=['Targets'])

# Check the info of the modified data
# print(iris_data.info())

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random_state=42)

# Initialize the Random Forest classifier
rf_classifier = RandomForestClassifier(n_estimators=100, random_state=42)

# Fit the classifier to the training data
rf_classifier.fit(X_train, y_train)

# Predict on the test data
y_pred = rf_classifier.predict(X_test)

# Evaluate the classifier
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")

# Print classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))

# Print confusion matrix
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

Results:

```
Accuracy: 0.98
Classification Report:
              precision    recall  f1-score   support

     0         1.00      1.00      1.00        23
     1         0.95      1.00      0.97        19
     2         1.00      0.94      0.97        18

   accuracy                0.98        60
  macro avg              0.98        0.98        0.98        60
weighted avg              0.98        0.98        0.98        60

Confusion Matrix:
[[23  0  0]
 [ 0 19  0]
 [ 0  1 17]]
```

11) Implement Boosting ensemble method on a given dataset.

Algorithm

Boosting ensemble method

Algorithm:

- 1) Import libraries
- 2) Load the dataset
- 3) Data preprocessing involves operations of features & dataset
- 4) Split dataset to train samples
- 5) Initialize adaboost classifier with specified no. of estimator & base estimator
- 6) Train model using training data
- 7) Make predictions for test sample using trained model
- 8) Evaluate model.

Output : model accuracy score : 0.833.

PAGE NO :

DATE :

Code:

```
In [4]: from sklearn.linear_model import LogisticRegression
        from sklearn.ensemble import AdaBoostClassifier
        from sklearn import metrics
        from sklearn import datasets

In [5]: import pandas as pd
        import matplotlib.pyplot as plt
        from sklearn.model_selection import train_test_split

In [6]: # Load the iris dataset
        iris = datasets.load_iris()
        X = pd.DataFrame(iris.data, columns=['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width'])
        y = pd.DataFrame(iris.target, columns=['Targets'])

In [7]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.4,random_state=42)

In [9]: mylogregmodel = LogisticRegression()

In [10]: adabc = AdaBoostClassifier(n_estimators = 150, base_estimator = mylogregmodel, learning_rate = 1)

In [11]: model = adabc.fit(X_train, y_train)

/usr/local/lib/python3.10/dist-packages/sklearn/utils/validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples, ), for example using ravel().
  y = column_or_1d(y, warn=True)
/usr/local/lib/python3.10/dist-packages/sklearn/ensemble/_base.py:166: FutureWarning: `base_estimator` was renamed to `estimator` in version 1.2 and will be removed in 1.4.
  warnings.warn(

In [12]: y_pred = model.predict(X_test)

In [13]: metrics.accuracy_score(y_test, y_pred)
```

Results:

```
Out[13]: 0.9833333333333333
```